

Abad Naseer

AI Infrastructure & Platform Engineer · GPU · Kubernetes · LLM Inference

+92 340-5059429 | abad.naseerfast@gmail.com | linkedin.com/in/abadnaseer | abad.falcoflow.com | Islamabad, PK (remote)

SUMMARY

I make self-hosted LLM inference cheap enough and reliable enough to run in production. Five years across cloud infrastructure and platform engineering, now owning the path from GPU cloud to served token: GPU workers, Kubernetes scheduling, vLLM, gateway routing and the automation that makes any of it replaceable. I came to this from the cost side — roughly 30% off one AWS bill at 1M+ users, and about 60% off infrastructure spend on a migration off AWS with zero customer-facing downtime. I also build what runs on top, which is less common in infrastructure engineers than it should be.

EXPERIENCE

Platform Engineer, AI Infrastructure

Jun. 2026 – Present

Ccript Agency

Remote

- Own the path from GPU cloud to served token for an inference platform: GPU workers, K3s control plane and scheduling, vLLM as the serving runtime, and a gateway routing per model in front of it
- Automated the full GPU worker lifecycle with Ansible — provision, join the cluster, verify the GPU is genuinely usable, synchronise firewall state, tear down and rebuild — so a worker is a thing you replace rather than repair
- Ran failure engineering against the platform rather than trusting that pods showing **Running** meant working: repeatedly destroyed and recreated GPU workers to prove unattended recovery, which surfaced stale firewall rules, node-local storage assumptions, a GPU device-plugin race, dynamic worker addressing, and verification logic competing with the workload for the GPU
- Building out observability and benchmarking across the inference path — Prometheus and Grafana over GPU, Kubernetes, runtime and gateway metrics, plus time-to-first-token, inter-token latency, throughput, GPU utilisation and KV-cache capacity — so model and configuration choices are made on measurement
- Separately, migrated a multi-brand franchise reporting platform off a third-party dashboard (Explo, serving 150+ outlets across the United States) onto a native reporting layer, then repaired the pipelines, KPI accuracy, filtering, location mappings and BigQuery reconciliation underneath it

DevOps Engineer

Oct. 2025 – Apr. 2026

Mediatiz Foundation

Islamabad, Pakistan

- Provisioned and managed AWS infrastructure (ECS/Fargate, RDS with ProxySQL and read replicas, ALB, CloudFront, Lambda, WAF, CloudWatch) with Terraform for an LMS and mobile app at 1M+ registered users and ~20–30K daily active sessions; held 99.9% uptime through auto-scaling and on-call incident response
- Cut monthly AWS spend by 30% by eliminating idle ECS tasks, fixing misconfigured Auto Scaling groups, removing orphaned EC2 instances, and working with the dev team to resolve slow database queries
- Deployed a self-hosted Django + Ollama (Gemma 2B) inference service on local infrastructure, removing a dependency on external model APIs and keeping student data on-premise
- Built GitHub Actions pipelines for Android, Laravel, Django and Next.js releases with automated ECS rolling deploys, replacing a manual release process that took 2–3 hours; ran OWASP testing against staging and updated WAF rules from the findings

DevOps Engineer

Dec. 2023 – Oct. 2025

Poshmaal Technologies

Islamabad, Pakistan

- Led the migration of firefly.online, an IoT SaaS platform, from AWS to self-managed infrastructure — EC2, RDS, S3, WAF and DNS — for roughly 60% lower infrastructure spend
- Ran Node.js and Spring Boot behind Nginx through a blue-green cutover, switching live traffic after three weeks of parallel validation with no customer-facing downtime
- Built dev, staging and production on bare metal; containerised the service estate with Docker and deployed to self-hosted Kubernetes via Helm with HPA, RBAC and resource limits per environment
- Built Jenkins and GitHub Actions pipelines with parallel stages and artifact caching, taking deployment time from ~35 to ~10 minutes; set up Prometheus, Grafana and Loki across all environments with per-service dashboards and multi-channel alerting

DevOps Consultant (Freelance)

Dec. 2021 – Nov. 2023

Fiverr & Upwork

Remote

- Delivered 20+ containerisation, CI/CD and server infrastructure projects for small-business clients; Docker-based deployments, GitHub Actions pipelines, Nginx configuration and SSL automation with Certbot, using Ansible for repeatable provisioning
- Deployed self-hosted Kubernetes for several clients with Ingress, HPA and RBAC, and wrote Bash tooling for routine server maintenance and patching

SELECTED PROJECTS

- Conversational commerce platform** | *FastAPI, pgvector, Gemini, Next.js* | smartzees.com 2026 – Present
- Three production AI assistants on one shared architecture — grocery ordering, home services and community documents — with the backend owning all state, search and money so the model only classifies intent and phrases the reply
 - Replaced per-request database scans with an in-memory vector index: catalog search fell from 6,456 ms to 8.2 ms and database reads per message from 25,701 to 2, with results identical to the old path on every query tested
 - Moved voice off the browser speech API onto server-side speech-to-text and synthesis, taking the product out of Chrome-only and giving the assistant one consistent voice
- Kubernetes GitOps platform** | *Kubernetes, ArgoCD, Helm, Prometheus, Traefik* 2025
- Deployed a 10-service polyglot application (Go, Java, Node.js, Python) on a self-hosted bare-metal cluster via Helm, with all cluster state in Git and ArgoCD auto-sync and self-healing on drift
 - CI never deploys: each pipeline builds and tags an image by commit SHA and commits a per-service tag override back to the GitOps repository, so one service rolls without waiting on the other nine
 - Deployed kube-prometheus-stack via ArgoCD tuned for an 8 GB host; wrote ServiceMonitors and instrumented cart and payment services with prom-client, verified end to end in Grafana via PromQL
- LogSense — log anomaly detection and RCA** | *Python, FastAPI, MLflow, Airflow, DVC* 2024
- Trained NLP models to classify infrastructure log anomalies and surface likely root causes, served through a FastAPI container with webhook-driven routing to Slack, Teams and PagerDuty
 - Tracked experiments with MLflow, versioned datasets with DVC and scheduled retraining with Airflow; cut alert noise by grouping correlated anomalies before paging on-call

EDUCATION

NUCES-FAST University Islamabad, Pakistan
Bachelor of Science in Software Engineering Sep. 2020 – Jun. 2025

TECHNICAL SKILLS

AI infrastructure: vLLM, GPU scheduling and node lifecycle, KV-cache reuse, gateway routing, inference benchmarking (TTFT, inter-token latency, throughput), Ollama, RAG, pgvector, HuggingFace

Containers & orchestration: Kubernetes, K3s, Helm, ArgoCD (GitOps), Docker (multi-stage), Nginx, Traefik, HPA, RBAC

Cloud: AWS (EC2, ECS/ECR, RDS, ALB, S3, CloudFront, WAF, Lambda, CloudWatch, Route 53, IAM, VPC), GCP (GKE, BigQuery, Cloud SQL), bare metal and self-managed

IaC & CI/CD: Terraform, Ansible, GitHub Actions, Jenkins, GitLab CI, Bash

Observability: Prometheus (ServiceMonitor, PromQL, prom-client), Grafana, Loki, Alertmanager, CloudWatch

Security: least-privilege IAM and RBAC, AWS WAF, SSL/TLS (Let's Encrypt), UFW/iptables, fail2ban, OWASP testing (Kali, SQLMap), SonarQube

Languages & data: Python (FastAPI, Django), Bash, TypeScript, Node.js, PostgreSQL, MySQL, Redis, Kafka

Pursuing: AWS Certified Solutions Architect – Associate (SAA-C03), Certified Kubernetes Administrator (CKA)